

Team Name: She+2 (Team\_30)

Architecture: MobileNetV3-Large with Dynamic Quantization.

Parameters: ~3 Million (Original is 5.4M, reduced via quantization).

Training Strategy: Transfer learning using ImageNet weights. Fine-tuned with data augmentation (RandomRotation, ColorJitter, HorizontalFlip) to ensure robustness against noise and lighting variations. Optimizer: Adam. Loss: CrossEntropy.

Ethical Considerations: The model is trained on [Gender Classification Dataset]. We acknowledge potential biases in facial recognition regarding ethnicity and age represented in the dataset.